

lab_gene_partial

May 22, 2019

1 Lab: PCA, LDA and Logistic Regression for Gene Expression Data

In this lab, we use logistic regression to predict biological characteristics ("phenotypes") from gene expression data. In addition to the concepts in breast cancer demo, you will learn to:

- Handle missing data
- Compute and visualize PCA and LDA coefficients
- Combine PCA and LDA with scaling
- Perform multi-class logistic classification on PCA and LDA outputs.
- Evaluate multi-class logistic classification with K-fold validation

1.1 Background

Genes are the basic unit in the DNA and encode blueprints for proteins. When proteins are synthesized from a gene, the gene is said to "express". Micro-arrays are devices that measure the expression levels of large numbers of genes in parallel. By finding correlations between expression levels and phenotypes, scientists can identify possible genetic markers for biological characteristics.

The data in this lab comes from:

<https://archive.ics.uci.edu/ml/datasets/Mice+Protein+Expression>

In this data, mice were characterized by three properties:

- Whether they had down's syndrome (trisomy) or not
- Whether they were stimulated to learn or not
- Whether they had a drug memantine or a saline control solution.

With these three choices, there are 8 possible classes for each mouse. For each mouse, the expression levels were measured across 77 genes. We will see if the characteristics can be predicted from the gene expression levels. This classification could reveal which genes are potentially involved in Down's syndrome and if drugs and learning have any noticeable effects.

1.2 Load the Data

We begin by loading the standard packages.

```
In [25]: # Imports
import numpy as np
import matplotlib.pyplot as plt
import matplotlib
```

```

matplotlib.rcParams.update({'font.size':16})
%matplotlib inline
import matplotlib.image as mpimg
from pylab import rcParams

import random
import math
from numpy.linalg import inv

import pandas as pd
from sklearn import linear_model, preprocessing

from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

```

Use the `pd.read_excel` command to read the data from https://archive.ics.uci.edu/ml/machine-learning-databases/00342/Data_Cortex_Nuclear.xls into a dataframe `df`. Use the `index_col` option to specify that column 0 is the index. Use the `df.head()` to print the first few rows.

```

In [26]: # Import the data
         # (Data from: https://archive.ics.uci.edu/ml/datasets/Mice+Protein+Expression)

io = 'https://archive.ics.uci.edu/ml/machine-learning-databases/00342/\
Data_Cortex_Nuclear.xls'
df = pd.read_excel(io, index_col=0)
df.head()

```

```

Out [26]:

```

	DYRK1A_N	ITSN1_N	BDNF_N	NR1_N	NR2A_N	pAKT_N	pBRAF_N	\
MouseID								
309_1	0.503644	0.747193	0.430175	2.816329	5.990152	0.218830	0.177565	
309_2	0.514617	0.689064	0.411770	2.789514	5.685038	0.211636	0.172817	
309_3	0.509183	0.730247	0.418309	2.687201	5.622059	0.209011	0.175722	
309_4	0.442107	0.617076	0.358626	2.466947	4.979503	0.222886	0.176463	
309_5	0.434940	0.617430	0.358802	2.365785	4.718679	0.213106	0.173627	

	pCAMKII_N	pCREB_N	pELK_N	...	pCFOS_N	SYP_N	H3AcK18_N	\
MouseID				...				
309_1	2.373744	0.232224	1.750936	...	0.108336	0.427099	0.114783	
309_2	2.292150	0.226972	1.596377	...	0.104315	0.441581	0.111974	
309_3	2.283337	0.230247	1.561316	...	0.106219	0.435777	0.111883	
309_4	2.152301	0.207004	1.595086	...	0.111262	0.391691	0.130405	
309_5	2.134014	0.192158	1.504230	...	0.110694	0.434154	0.118481	

	EGR1_N	H3MeK4_N	CaNA_N	Genotype	Treatment	Behavior	class
MouseID							
309_1	0.131790	0.128186	1.675652	Control	Memantine	C/S	c-CS-m

309_2	0.135103	0.131119	1.743610	Control	Memantine	C/S	c-CS-m
309_3	0.133362	0.127431	1.926427	Control	Memantine	C/S	c-CS-m
309_4	0.147444	0.146901	1.700563	Control	Memantine	C/S	c-CS-m
309_5	0.140314	0.148380	1.839730	Control	Memantine	C/S	c-CS-m

[5 rows x 81 columns]

This data has missing values. The site:

http://pandas.pydata.org/pandas-docs/stable/missing_data.html

has an excellent summary of methods to deal with missing values. Following the techniques there, create a new data frame df1 where the missing values in each column are filled with the mean values from the non-missing values.

```
In [27]: # pd.set_option('display.max_rows', -1)
# pd.reset_option(pat='display')

col = df.columns # Column labels
xnames = np.array(col[0:-4]) # Array of gene names
onames = np.array(col[-4:]) # Array of object-type column names
dfa = df[xnames]
dfb = df[onames]
# New data frame with missing values filled by linear interpolation
df0 = dfa.astype(float).interpolate(axis=0)
# NaNs in first rows weren't filled
df0 = df0.astype(float).interpolate(axis=0, limit_direction='backward')
# print(df0)
# print(np.isnan(df0).any())
```

We next get the data as numpy arrays. For the predictors, X, we will use the expression levels of the ngene=77 genes. The expression levels are stored in the first 77 columns of the dataframe df1.

- Set xnames = the names of genes (you can get them from df1.columns)
- Set X = a numpy array with the values of the expression levels. (you can get this from df1[xnames].values)

```
In [28]: ngene = 77 # number of genes
df1 = pd.concat([df0, dfb], axis=1) # New data frame
X = df1[xnames].values # Array of the values of the expression levels
```

Now run the following code which will extract the class of each measurement into a vector y. The values y will have values 0 to 7 corresponding to the 8 classes. Our goal will be to predict y from X.

```
In [29]: # Extract the class of each measurement into a vector y
ystr = df1['class'].values
vals, y = np.unique(ystr, return_inverse=True)
```

Next, split the data into training and test. You can use the train_test_split function. Set shuffle=True and test_size=0.5.

```
In [30]: # Split the data into training and test sets
from sklearn.model_selection import train_test_split
Xtr, Xts, ytr, yts = train_test_split(X, y, shuffle=True, test_size=0.5)
```

1.3 PCA on the Data

We will first try to perform PCA. With PCA, it is important to first scale the data matrix to remove the mean and normalize the features by their variance. We can do the scaling and PCA in two steps using routines from sklearn package:

- Create a scaling object, `scaler = StandardScaler(...)` and fit and transform the scaler on the training data, `Xtr`.
- Create a PCA object, `pca = PCA(...)` and fit the PCA coefficients on the scaled data. In order that we can visualize the results, set `n_components=2`.

```
In [31]: # Does the contain contain an NaN value?
```

```
print(np.isnan(Xtr).any())  
print(np.isnan(ytr).any())  
print(np.isnan(X).any())
```

```
False
```

```
False
```

```
False
```

```
In [32]: from sklearn.decomposition import PCA  
        from sklearn.preprocessing import StandardScaler
```

```
# Standard scaling function  
def f_StandardScale(Xtr,Xts):  
    scaler = StandardScaler()  
    Xtr_ = scaler.fit_transform(Xtr)  
    Xts_ = scaler.transform(Xts)  
    return Xtr_,Xts_
```

Now use the transform method to transform the test data `Xts` through the scaler and `pca` objects. Create a scatter plot using the `plt.scatter` function of the points from the two classes. Use different colors for each class. You will see that the PCA representation does not differentiate the classes well along the two components

```
In [33]: # Create a scaling object, and fit and transform
```

```
# the scaler on the training data.  
[Xtr_,Xts_] = f_StandardScale(Xtr,Xts)
```

```
# Create a PCA object and fit the  
# PCA coefficients on the scaled data.  
pca = PCA(n_components=2)  
Xtr_pca = pca.fit_transform(Xtr_)  
# Transform the test data through the scalar and PCA objects  
Xts_pca = pca.transform(Xts_)
```

```
In [35]: # Scatter plot of the data
```

```

num_classes = 8
fig, ax = plt.subplots()
cdict = ['k', 'r', 'orange', 'y', 'g', 'c', 'b', 'm'] # Color dictionary


datax = Xts_pca.T[0]
datay = Xts_pca.T[1]

for cc in np.arange(num_classes):
    ivec = np.where(yts == cc)
    # Generate 'num_classes' tuples equally spread
    # in hue space, then convert to RGB.
    # color = np.array(colorsys.hsv_to_rgb(cc*1.0/num_classes, 0.5, 0.5))
    # (doesn't work too well)
    color = cdict[cc]
    ax.scatter(datax[ivec], datay[ivec], c=color, label=cc)

plt.xlabel('x1')
plt.ylabel('x2')

ax.grid()
ax.legend(loc='best', bbox_to_anchor=(1,1))
rcParams['figure.figsize'] = 16, 8
plt.show()

```



1.4 LDA on the Data

A better way to transform data in a way that separates classes is LDA. The `sklearn` has excellent routines for LDA.

- As in the PCA case, create a scaling object, `scaler = StandardScaler(...)` and fit and transform the scaler on the training data, `Xtr`.
- Next create an LDA object, `lda = LinearDiscriminantAnalysis(...)`. To avoid ill-conditioning, set `shrinkage='auto'` and `solver='eigen'`. Fit the LDA transform from the scaled output of scaler.

```
In [36]: # Create a scaling object, and fit and transform
# the scaler on the training data.
[Xtr_,Xts_] = f_StandardScale(Xtr,Xts)

# Create an LDA object and fit the LDA transform.
lda = LinearDiscriminantAnalysis(n_components=2, \
                                shrinkage='auto', solver='eigen')
Xtr_lda = lda.fit_transform(Xtr_, ytr)
# Transform the test data through the scalar and LDA objects
Xts_lda = lda.transform(Xts_)

# print(Xts_lda)
```

Now transform the test data `Xts` through the scaler and `lda` objects. Create a scatter plot using the `plt.scatter` function of the points from the two classes. Use different colors for each class. You will see that LDA results in much better separation.

```
In [37]: # Scatter plot of the data


num_classes = 8
fig, ax = plt.subplots()
cdict = ['k', 'r', 'orange', 'y', 'g', 'c', 'b', 'm'] # Color dictionary

datax = Xts_lda.T[0]
datay = Xts_lda.T[1]

for cc in np.arange(num_classes):
    ivec = np.where(yts == cc)
    color = cdict[cc]
    ax.scatter(datax[ivec],datay[ivec],c=color,label=cc)

plt.xlabel('x1')
plt.ylabel('x2')

ax.grid()
ax.legend(loc='best', bbox_to_anchor=(1,1))
rcParams['figure.figsize'] = 16, 8
plt.show()
```



1.5 Logistic Regression on the LDA Data

We will now build a linear classifier from the LDA outputs. To fit the classifier, we use a three step pipeline:

- As above, create a scaling object, `scaler = StandardScaler(...)` and fit and transform the scaler on the training data, `Xtr`.
- Also, as above create an LDA object, `lda = LinearDiscriminantAnalysis(...)` and fit and transform the scaled training data. Call the transformed output `Ztr`.
- Create a logistic regression object, `logreg = linear_model.LogisticRegression(...)`. Set `solver='lbfgs'`, and `multi_class='auto'`. Fit the model on the transformed training data.

```
In [38]: # Create a scaling object, and fit and transform the scaler on the data.
[Xtr_,Xts_] = f_StandardScale(Xtr,Xts)

# Create an LDA object and fit the LDA transform.
lda = LinearDiscriminantAnalysis(n_components=8, \
                                shrinkage='auto', solver='eigen')
Ztr = lda.fit_transform(Xtr_, ytr)
# Transform the test data through the scalar and LDA objects
Xts_lda = lda.transform(Xts_)

# Create a logistic regression object
logreg = linear_model.LogisticRegression(solver='lbfgs', \
                                          multi_class = 'multinomial')

Xtr_log = logreg.fit(Ztr,ytr)
```

Now test the model on the test data:

- Scale the test data `Xts` with the `scaler.transform()` method
- Transform the scaled test data with the `lda.transform()` method
- Predict the class labels from `logreg.predict`. Call the outputs `yhat`.
- Measure the accuracy by comparing the outputs `yhat` with `yts`.

If you did everything correctly, you should get an accuracy of around 94%.

```
In [39]: # Predict the class labels
yhat = logreg.predict(Xts_lda)
ydiff = np.absolute(yhat-yts)
error = np.count_nonzero(ydiff)/np.size(ydiff)
per_acc = (1-error)*100
print('Accuracy:', round(per_acc,2), '%')
```

Accuracy: 96.67 %

1.6 K-Fold Cross Validation

K-Fold validation can yield better assessments of the accuracy when the training data is limited. Complete the following code to perform 5 fold validation.

```
In [21]: # A 5-fold cross validation of the data
from sklearn.model_selection import KFold
from sklearn.metrics import precision_recall_fscore_support
nfold = 5
kf = KFold(n_splits=nfold, shuffle=True)
acc = np.zeros(nfold)

for i, I in enumerate(kf.split(X)):
    # Get training and test data
    Itr, Its = I
    Xtr = X[Itr,:]
    ytr = y[Itr]
    Xts = X[Its,:]
    yts = y[Its]

    # Train the scaler, LDA, and logistic regression
    # model on the training data
    [Xtr_, Xts_] = f_StandardScale(Xtr, Xts)
    lda = LinearDiscriminantAnalysis(n_components=8, \
                                     shrinkage='auto', solver='eigen')
    Ztr = lda.fit_transform(Xtr_, ytr)
    # Transform the test data through the scalar and LDA objects
    Xts_lda = lda.transform(Xts_)
    logreg = linear_model.LogisticRegression(solver='lbfgs', \
                                              multi_class = 'multinomial')
    Xtr_log = logreg.fit(Ztr, ytr)
```



```

# Test the scaler, LDA, and regression model on the test data
yhat = logreg.predict(Xts_lda)

# Measure accuracy and store in acc[i]
ydiff = np.absolute(yhat-yts)
error = np.count_nonzero(ydiff)/np.size(ydiff)
per_acc = (1-error)*100
acc[i] = per_acc

# Print the mean and SE of the accuracy
acc_mean = np.mean(acc)
acc_std = np.std(acc)
print('Mean Accuracy:', round(acc_mean,2),'%')
print('Accuracy SE:', round(acc_std,2),'%')

/Users/peterracioppo/anaconda3/lib/python3.6/site-packages/sklearn/discriminant_analysis.py:44:
UserWarning)

```

Mean Accuracy: 97.31 %
Accuracy SE: 0.9 %

1.7 More Fun

Statistical analysis of genetic analysis is a rich area and there are several simple things that you can explore as a continuation of this lab:

- Larger datasets
- Combining the K-fold validation with parameter optimization
- Using sparse LDA or sparse regression

1.7.1 K-Fold Cross Validation + Parameter Optimization

```

In [23]: # Vary the number of folds
from sklearn.model_selection import KFold
from sklearn.metrics import precision_recall_fscore_support

fold_vec = np.arange(9)+2
for nn in fold_vec:
    nfold = nn
    kf = KFold(n_splits=nfold,shuffle=True)
    acc = np.zeros(nfold)

    for i, I in enumerate(kf.split(X)):
        # Get training and test data
        Itr, Its = I
        Xtr = X[Itr,:]
        ytr = y[Itr]

```

```

Xts = X[Its,:]
yts = y[Its]

# Train the scaler, LDA, and logistic
# regression model on the training data
[Xtr_,Xts_] = f_StandardScale(Xtr,Xts)
lda = LinearDiscriminantAnalysis(n_components=8, \
                                shrinkage='auto', solver='eigen')
Ztr = lda.fit_transform(Xtr_, ytr)
# Transform the test data through the scalar and LDA objects
Xts_lda = lda.transform(Xts_)
logreg = linear_model.LogisticRegression(solver='lbfgs', \
                                         multi_class = 'multinomial')

Xtr_log = logreg.fit(Ztr,ytr)

# Test the scaler, LDA, and regression model on the test data
yhat = logreg.predict(Xts_lda)

# Measure accuracy and store in acc[i]
ydiff = np.absolute(yhat-yts)
error = np.count_nonzero(ydiff)/np.size(ydiff)
per_acc = (1-error)*100
acc[i] = per_acc

# Print the mean and SE of the accuracy
acc_mean = np.mean(acc)
acc_std = np.std(acc)
print('Number of Folds:', nn)
print('Mean Accuracy:', round(acc_mean,2), '%')

```

```

/Users/peterracioppo/anaconda3/lib/python3.6/site-packages/sklearn/discriminant_analysis.py:44:
UserWarning)

```

```

Number of Folds: 2
Mean Accuracy: 96.11 %
Number of Folds: 3
Mean Accuracy: 96.57 %
Number of Folds: 4
Mean Accuracy: 96.67 %
Number of Folds: 5
Mean Accuracy: 97.41 %
Number of Folds: 6
Mean Accuracy: 97.78 %
Number of Folds: 7
Mean Accuracy: 97.69 %
Number of Folds: 8
Mean Accuracy: 97.87 %

```

Number of Folds: 9
Mean Accuracy: 97.5 %
Number of Folds: 10
Mean Accuracy: 97.59 %

1.7.2 Sparse LDA

```
In [24]: # Vary the number of LDA components
# A 5-fold cross validation of the data
from sklearn.model_selection import KFold
from sklearn.metrics import precision_recall_fscore_support
nfold = 5
kf = KFold(n_splits=nfold,shuffle=True)
acc = np.zeros(nfold)

comp_vec = np.arange(10)+1
for nc in comp_vec:

    for i, I in enumerate(kf.split(X)):
        # Get training and test data
        Itr, Its = I
        Xtr = X[Itr,:]
        ytr = y[Itr]
        Xts = X[Its,:]
        yts = y[Its]

        # Train the scaler, LDA, and logistic
        # regression model on the training data
        [Xtr_,Xts_] = f_StandardScale(Xtr,Xts)
        lda = LinearDiscriminantAnalysis(n_components=nc, \
                                         shrinkage='auto', solver='eigen')
        Ztr = lda.fit_transform(Xtr_, ytr)
        # Transform the test data through the scalar and LDA objects
        Xts_lda = lda.transform(Xts_)
        logreg = linear_model.LogisticRegression(solver='lbfgs', \
                                                  multi_class = 'multinomial')
        Xtr_log = logreg.fit(Ztr,ytr)

        # Test the scaler, LDA, and regression model on the test data
        yhat = logreg.predict(Xts_lda)

        # Measure accuracy and store in acc[i]
        ydiff = np.absolute(yhat-yts)
        error = np.count_nonzero(ydiff)/np.size(ydiff)
        per_acc = (1-error)*100
        acc[i] = per_acc
```

```

# Print the mean and SE of the accuracy
acc_mean = np.mean(acc)
acc_std = np.std(acc)
print('Number of LDA Components:', nc)
print('Mean Accuracy:', round(acc_mean,2), '%')

```

```

/Users/peterracioppo/anaconda3/lib/python3.6/site-packages/sklearn/discriminant_analysis.py:44:
UserWarning)

```

```

Number of LDA Components: 1
Mean Accuracy: 50.28 %
Number of LDA Components: 2
Mean Accuracy: 82.5 %
Number of LDA Components: 3
Mean Accuracy: 91.11 %
Number of LDA Components: 4
Mean Accuracy: 92.96 %
Number of LDA Components: 5
Mean Accuracy: 94.07 %
Number of LDA Components: 6
Mean Accuracy: 95.83 %
Number of LDA Components: 7
Mean Accuracy: 96.85 %
Number of LDA Components: 8
Mean Accuracy: 97.22 %
Number of LDA Components: 9
Mean Accuracy: 97.41 %
Number of LDA Components: 10
Mean Accuracy: 97.31 %

```